

ADT COSETTE

Cristian Ruiz

MADYNES Team

May 31, 2016 – young engineers seminar



Validation in Computer Science

Two classical approaches for validation:

- **Formal**: equations, proofs, etc.
- **Experimental**, on a scientific instrument.

In reality

Given the size and complexity of distributed systems, it is difficult to carry out complete analytic studies. **So, the experimental approach is commonly seen.**

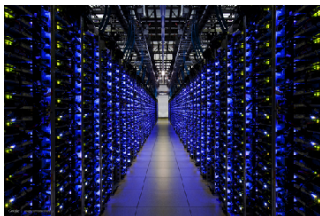
Need of

Reproducible research

A large-scale, shared testbed supporting high-quality, reproducible research on distributed systems:

- Highly configurable
- High Performance Computing, Grids, Peer-to-peer systems, Cloud computing

Another examples of testbeds: Chameleon, Cloudlab



Goal of the ADT COSETTE

Conceive, consolidate and extend a set of tools aimed at experimenting with distributed systems (Cloud, Grid, HPC, P2P)

Tasks

- Development of Ruby-Cute, a library that gathers useful procedures for experimenting with distributed systems
- **Extend Distem to meet Cloud and HPC research requirements**

Supervised by

Lucas Nussbaum, Emmanuel Jeanvoine

Outline

- ① Distem
- ② Evaluating HPC runtimes with Distem
- ③ Distem validation for HPC applications
- ④ Conclusions

An emulator for distributed systems

Take your **real application**

Run it on a **cluster**

And use **Distem** to **alter the platform**
so it **matches the experimental conditions you need**



+



Heterogeneous nodes
Long distance networks
Faults, perf. variations
Grid, Cloud, P2P features
...

Distem features

The features of Distem include:

- running many virtual nodes on each physical node
- emulation of CPU performance, network topologies, I/O speed

Distem uses modern Linux functionality:

- Linux containers
- control groups
- CPU frequency scaling
- traffic control
- I/O throttling

Outline

- ① Distem
- ② Evaluating HPC runtimes with Distem
- ③ Distem validation for HPC applications
- ④ Conclusions

HPC runtimes

- According to the IESP report a strong effort must be made on improving HPC software stacks
- One of the main parts of this stack is dedicated to **HPC runtime**
- HPC runtime enables the execution, managing and debugging of parallel applications
- OpenMPI, Charm++, CUDA, etc.

For this work we focus on studying HPC runtimes

Dongarra, Jack et Al., *The International Exascale Software Project Roadmap*,
International Journal of High Performance Computer Applications, 2011

Evaluating current HPC runtimes

Several properties to evaluate

- Programmability
- Scalability
- Fault tolerance
- Load balancing

We focus on

- Fault tolerance: more components \Rightarrow shorter MTBF (Mean Time Between Failures)
- Load balancing: Cloud computing, Green computing, Data centers' policies

Evaluating current HPC runtimes

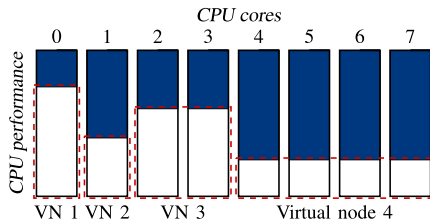
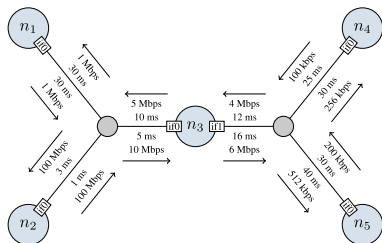
- Carrying out evaluation under complex realistic conditions is **hard**
- Simulator:
 - simplified assumptions 😞
 - lower realism 😞
 - not possible to run a complete software stack 😞
- Real platform:
 - expensive 😞
 - lacks of reproducibility 😞

In this task

We integrated the following improvements in order to make possible the evaluation of HPC runtimes:

- Evolving experimental conditions
- Failure injection framework
- Event injection framework

Evolving experimental conditions



- Heterogeneous conditions can be created: CPU frequencies, different IO and network capabilities
- These features can be updated dynamically
- This is useful to achieve complex experiments where the platform is modified, like it could happened in reality

Failure injection framework

- We take into account failures that provoke a loss of the node (very common failures)
- Nodes can be lost in three different ways:
 - **Graceful**: the node is shut down cleanly, using an operating system command
 - **Soft**: the node is forced to shut down
 - **Hard**: the node failed abruptly
- We do not take into account byzantine failures

- Increase the reproducibility of experiments
- Distem supports the following modifications for a given set of nodes:
 - CPU frequency
 - Network capabilities (latency and bandwidth)
 - Failures
- These modifications can be injected using a deterministic behavior or using a probabilistic distribution

Experiment setup

- We evaluate Charm++, OpenMPI and MPICH runtimes
- Charm++: Jacobi3D and Stencil3D
- MPI-based runtimes: NAS parallel benchmarks
- 3 Grid'5000 clusters located in two sites
- Experimental evaluation:
 - *Failure detection of HPC runtimes*
 - *Evaluation of load balancing strategies in Charm++*
 - *Validity of fault injection mechanism*

Failure detection of HPC runtimes

- We run an application on top of the HPC runtime
- We inject different types of faults and observe how the HPC runtime reacts

Failure	Runtime					
	Charm++		OpenMPI		MPICH	
	Detected	Action	Detected	Action	Detected	Action
Graceful	Yes	C	Yes	H	Yes	E
Soft	Yes	C	Yes	H	Yes	E
Hard	No	-	Yes	H	Yes	E

Table: Failure detection. C refers to the roll-back of the application to the previous checkpoint, H refers to the fact that processes hang, E refers to the termination of MPI processes

Evaluating load balancing strategies in Charm++

- We create a platform composed 128 vnodes distributed over 8 physical nodes.
- We experiment with two different scenarios:
 - **Heterogeneous**: half of the vnodes have a CPU clock reduced to 50 %
 - **Dynamic**: the available CPU power of a sub-part of the vnodes is dynamic.

The event injection framework was used to automate the creation of these scenarios

Evaluating load balancing strategies in Charm++

Running Stencil3D using 128 processes in the heterogeneous platform

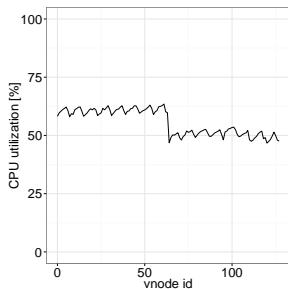


Figure: LBOff
Walltime: 341 secs

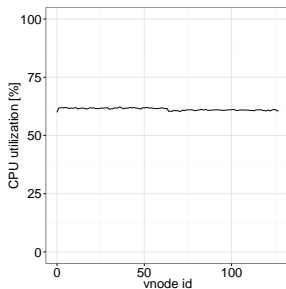


Figure: RefineLB
Walltime: 320 secs

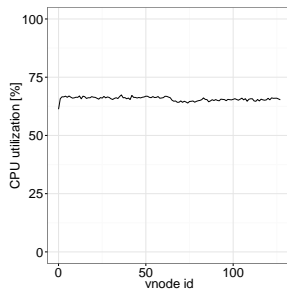


Figure: Hybrid
Walltime: 356 secs

Evaluating load balancing strategies in Charm++

Running Stencil3D using 128 processes in the dynamic platform

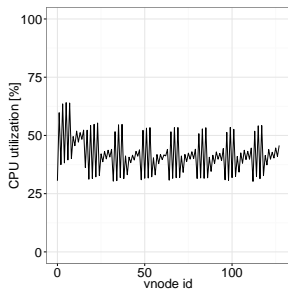


Figure: LBOff
Walltime: 347 secs

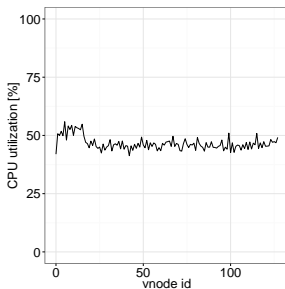


Figure: RefineLB
Walltime: 322 secs

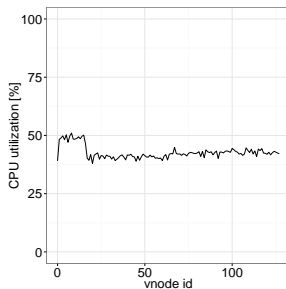


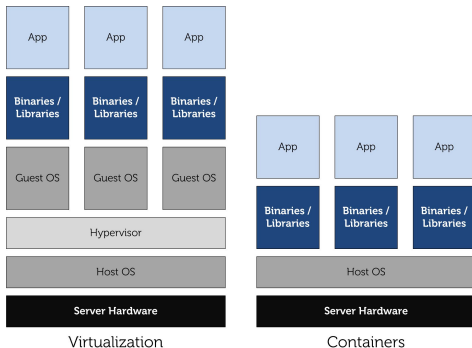
Figure: Hybrid
Walltime: 359 secs

Outline

- ① Distem
- ② Evaluating HPC runtimes with Distem
- ③ Distem validation for HPC applications
- ④ Conclusions

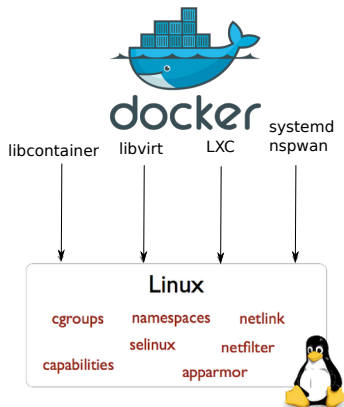
Containers

Containers refers generally to **Operating-system-level virtualization**, where the **kernel** of an operating system allows for multiple isolated **user-space** instances.



namespaces and cgroups

- Both features incorporated in Linux kernel since 2006 (Linux 2.6.24)
- Several container solutions: LXC, libvirt, libcontainer, systemd-nspawn, Docker



In this task, we answer:

- What is the overhead of oversubscription using different versions of Linux kernel?
- What is the impact of running an HPC workload with several MPI processes inside containers?
- What is the impact of network technology ?

Experimental setup

Hardware

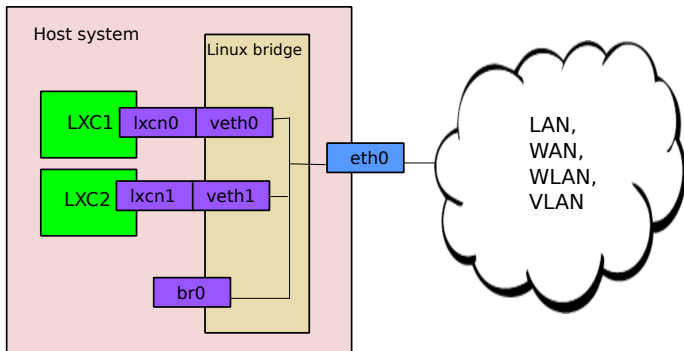
- Cluster in Grid'5000 Testbed where each node is equipped with two Intel Xeon E5-2630v3 processors (with 8 cores each), 128 GB of RAM and a 10 GbE adapter
- Our experimental setup included up to 64 machines

Software

- Debian Jessie, Linux kernel versions: 3.2, 3.16 and 4.0, OpenMPI and NPB. We instrumented the benchmarks: LU, EP, CG, MG, FT, IS using TAU

Network setup

- Veth pair + Linux bridge
- Veth pair + OpenvSwitch
- MACVLAN or SR-IOV
- Phys

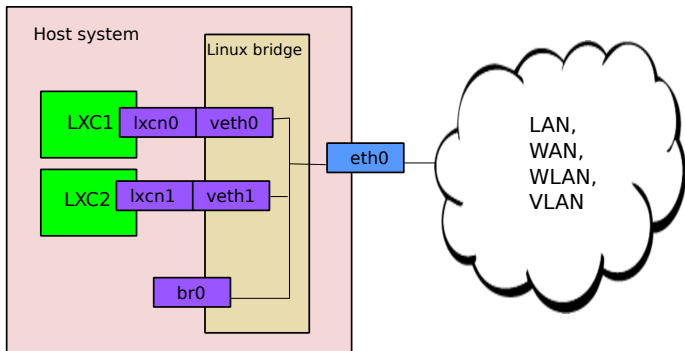


What did we find?

- There is important performance degradation provoked by veth for Linux kernels < 3.11. Running 2 containers per physical machine:
 - 3.2: 1577.78%
 - 3.16: 22.67%
 - 4.0: 2.40%
- Overhead present in MPI communication
- Since Linux kernel version 3.11, TSO was enabled in veth
- Memory bound applications and application that use all to all MPI communication are the most affected by oversubscription
- Performance issues can appear only at certain scale (e.g. 180% overhead with 128 MPI processes for CG benchmark).

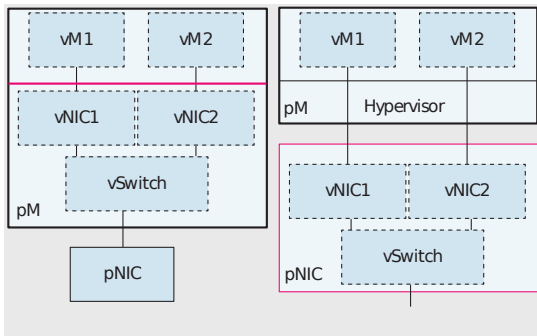
Interconnection comparison

- Veth pair + Linux bridge
- Veth pair + OpenvSwitch
- SR-IOV
- Using Linux kernel 4.3



SR-IOV explained

- Veth pair + Linux bridge
- Veth pair + OpenvSwitch
- SR-IOV

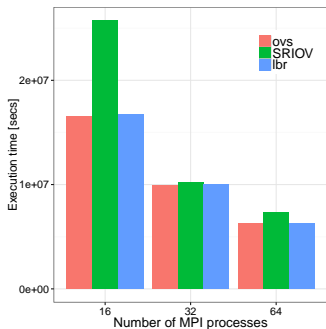


Network Virtualization and Software Defined Networking for Cloud Computing: A survey
Raj Jain and Subharthi Paul, Washington University

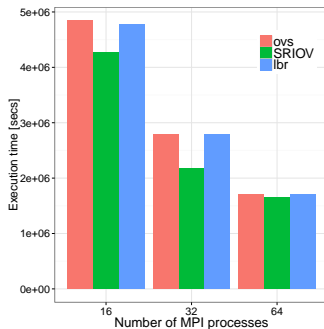
Multi-machine (4 nodes)

- 4 containers per machine
- Each container configured with 4 cores

(a) LU Class B



(b) EP Class B



What did we find?

- Using one machine there is a significant overhead of SR-IOV
- Multi-machine setups results varies from benchmark to benchmark there is no significant improvement of using SR-IOV
- Communication of container in the same machine are really affected using SR-IOV
- OpenVSwitch presents the best performance

Outline

- ① Distem
- ② Evaluating HPC runtimes with Distem
- ③ Distem validation for HPC applications
- ④ Conclusions

- Being able to execute experiments on a large set of platform configurations in a repeatable way is a sound basis to design and improve the HPC runtimes in the future
- **Distem:**
 - offers realistic experimental conditions
 - simplified the uncovering of problems in the failure handling for widely used HPC runtimes
 - enables experimenters to easily simulate perturbations and heterogeneity of nodes

What did we find?

- There is important performance degradation provoked by **veth** for Linux kernels < 3.11
- Container placing plays an important role under oversubscription
- Memory bound applications and application that use **all to all MPI** communication are the most affected by oversubscription
- More details about the results:

Cristian Ruiz et Al., *Performance Evaluation of Containers for HPC*, 11th Workshop on Virtualization in High-Performance Cloud Computing, Vienna, Austria, 2015